

Application Acceleration with High-Level-Synthesis

Dec. 29, 2022

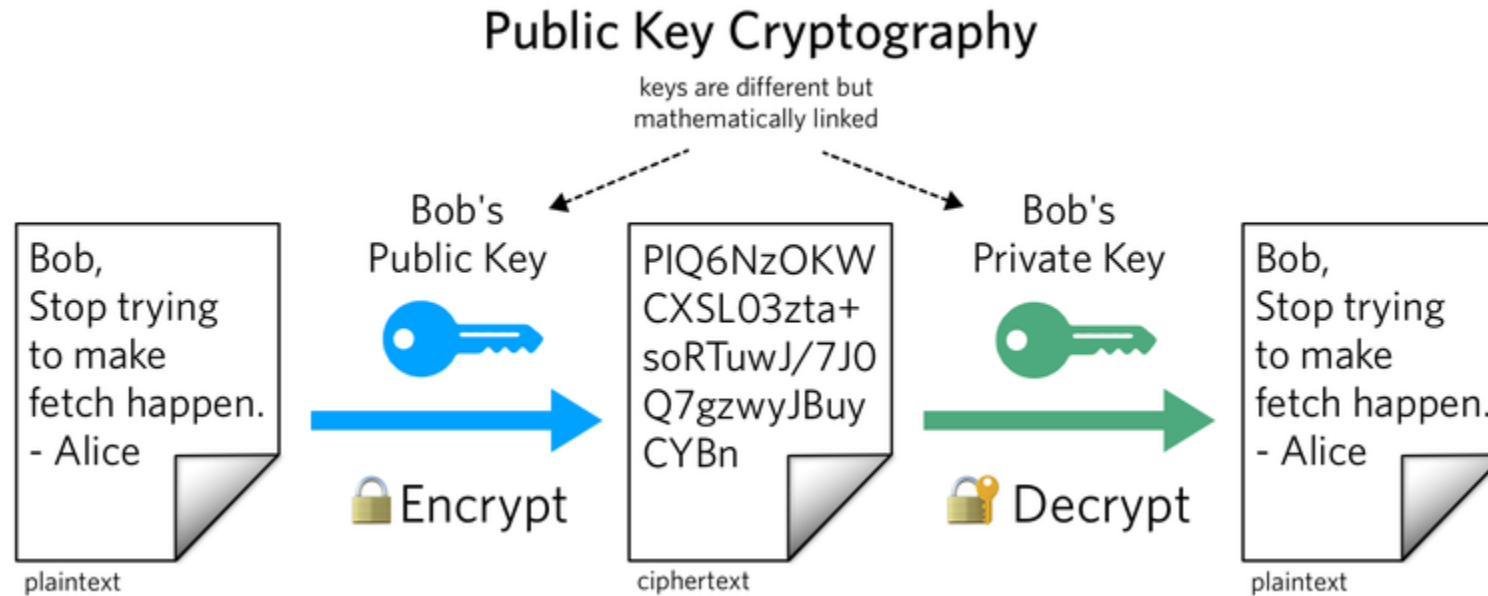
HLS Final Project - Dilithium

Presenter: Liang-Hsin Lin, Yu-Cheng Lin

Graduate Institute of Electronics Engineering, National Taiwan University



Cryptography




- **Encrypting with a public key**
 - Sending messages only the intended recipient can read
- **Signing with your private key**
 - Verifying that you're the one who sent a message

Post-Quantum Cryptography

- Current popular algorithms could be easily solved on a sufficiently powerful quantum computer running Shor's algorithm
- Post-quantum cryptography is secure against a cryptanalytic attack by a quantum computer
- Computation complexity is high
 - Need efficient algorithm and dedicated hardware codesign
- Hardware acceleration design reference in the past few years
 - ISSCC'22, VLSI'22, ISSCC'23
 - ➔ **Fast architecture search using HLS**

PQC Algorithm: Dilithium

- Basic component:
 - Polynomial ring
 - Degree of 256



Gen

```

01  $\mathbf{A} \leftarrow R_q^{k \times \ell}$ 
02  $(\mathbf{s}_1, \mathbf{s}_2) \leftarrow S_\eta^\ell \times S_\eta^k$ 
03  $\mathbf{t} := \mathbf{A}\mathbf{s}_1 + \mathbf{s}_2$ 
04 return  $(pk = (\mathbf{A}, \mathbf{t}), sk = (\mathbf{A}, \mathbf{t}, \mathbf{s}_1, \mathbf{s}_2))$ 
    
```

Sign(sk, M)

```

05  $\mathbf{z} := \perp$ 
06 while  $\mathbf{z} = \perp$  do
07    $\mathbf{y} \leftarrow S_{\gamma_1 - 1}^\ell$ 
08    $\mathbf{w}_1 := \text{HighBits}(\mathbf{A}\mathbf{y}, 2\gamma_2)$ 
09    $c \in B_\tau := H(M \parallel \mathbf{w}_1)$ 
10    $\mathbf{z} := \mathbf{y} + c\mathbf{s}_1$ 
11   if  $\|\mathbf{z}\|_\infty \geq \gamma_1 - \beta$  or  $\|\text{LowBits}(\mathbf{A}\mathbf{y} - c\mathbf{s}_2, 2\gamma_2)\|_\infty \geq \gamma_2 - \beta$ , then  $\mathbf{z} := \perp$ 
12 return  $\sigma = (\mathbf{z}, c)$ 
    
```

Verify($pk, M, \sigma = (\mathbf{z}, c)$)

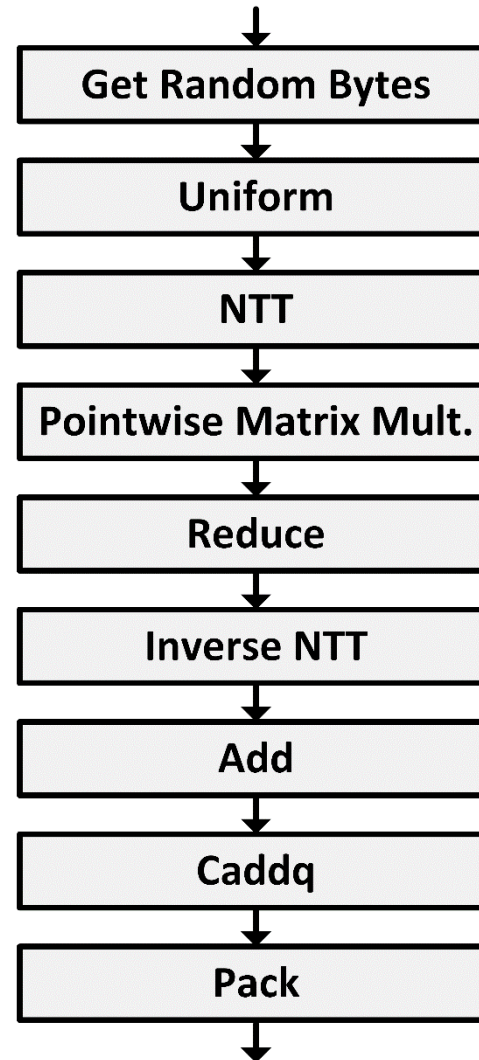
```

13  $\mathbf{w}'_1 := \text{HighBits}(\mathbf{A}\mathbf{z} - c\mathbf{t}, 2\gamma_2)$ 
14 if return  $\llbracket \|\mathbf{z}\|_\infty < \gamma_1 - \beta \rrbracket$  and  $\llbracket c = H(M \parallel \mathbf{w}'_1) \rrbracket$ 
    
```

Algorithm Workflow

Key generation for example

- Uniform Sampling
- Element-wise Operations
 - NTT/Inverse NTT
 - Addition/Subtraction
 - Reduce
 - Caddq
- Pack



```
uint8_t tr[SEEDBYTES];
const uint8_t *rho, *rho_prime, *key;
polyvec1 mat[K];
polyvec1 s1, s1hat;
polyveck s2, t1, t0;

/* Get randomness for rho, rho_prime and key */
randombytes(seedbuf, SEEDBYTES);
shake256(seedbuf, 2*SEEDBYTES + CRHBYTES, seedbuf, SEEDBYTES);
rho = seedbuf;
rho_prime = rho + SEEDBYTES;
key = rho_prime + CRHBYTES;

/* Expand matrix */
polyvec_matrix_expand(mat, rho);

/* Sample short vectors s1 and s2 */
polyvec1_uniform_eta(&s1, rho_prime, 0);
polyveck_uniform_eta(&s2, rho_prime, L);

/* Matrix-vector multiplication */
s1hat = s1;
polyvec1_ntt(&s1hat);
polyvec_matrix_pointwise_montgomery(&t1, mat, &s1hat);
polyveck_reduce(&t1);
polyveck_invntt_tomont(&t1);

/* Add error vector s2 */
polyveck_add(&t1, &t1, &s2);


/* Extract t1 and write public key */
polyveck_caddq(&t1);
polyveck_power2round(&t1, &t0, &t1);
pack_pk(pk, rho, &t1);

/* Compute H(rho, t1) and write secret key */
shake256(tr, SEEDBYTES, pk, CRYPTO_PUBLICKEYBYTES);
pack_sk(sk, rho, tr, key, &t0, &s1, &s2);
```

Algorithm Analysis – CPU Profiling

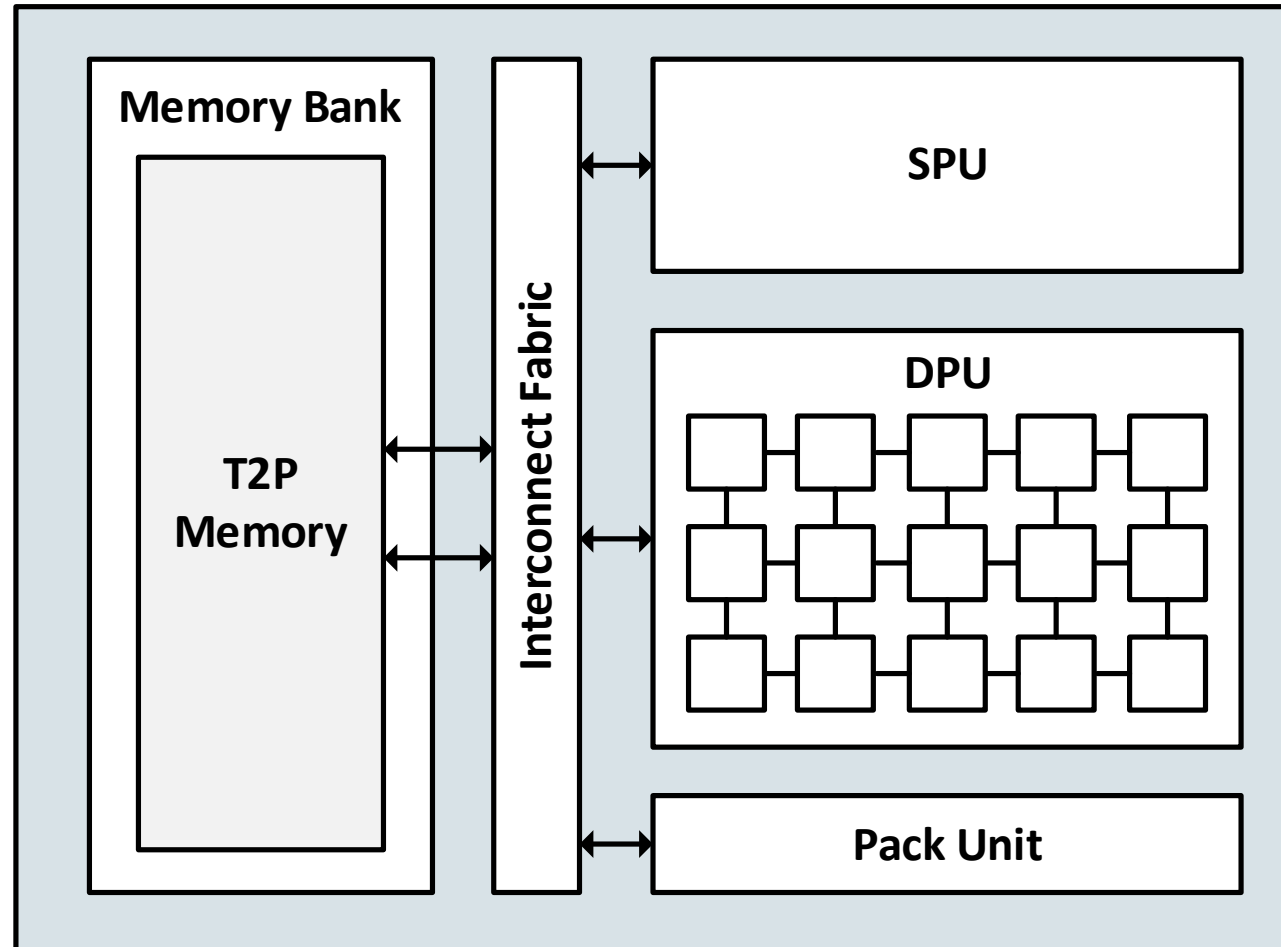
- Run on Intel Core i7-9700
- Computation bottleneck:
 - Uniform sampling
 - Matrix multiplication
 - Number theoretic transform (NTT)
 - Inverse NTT

➔ accounts for 82% execution time

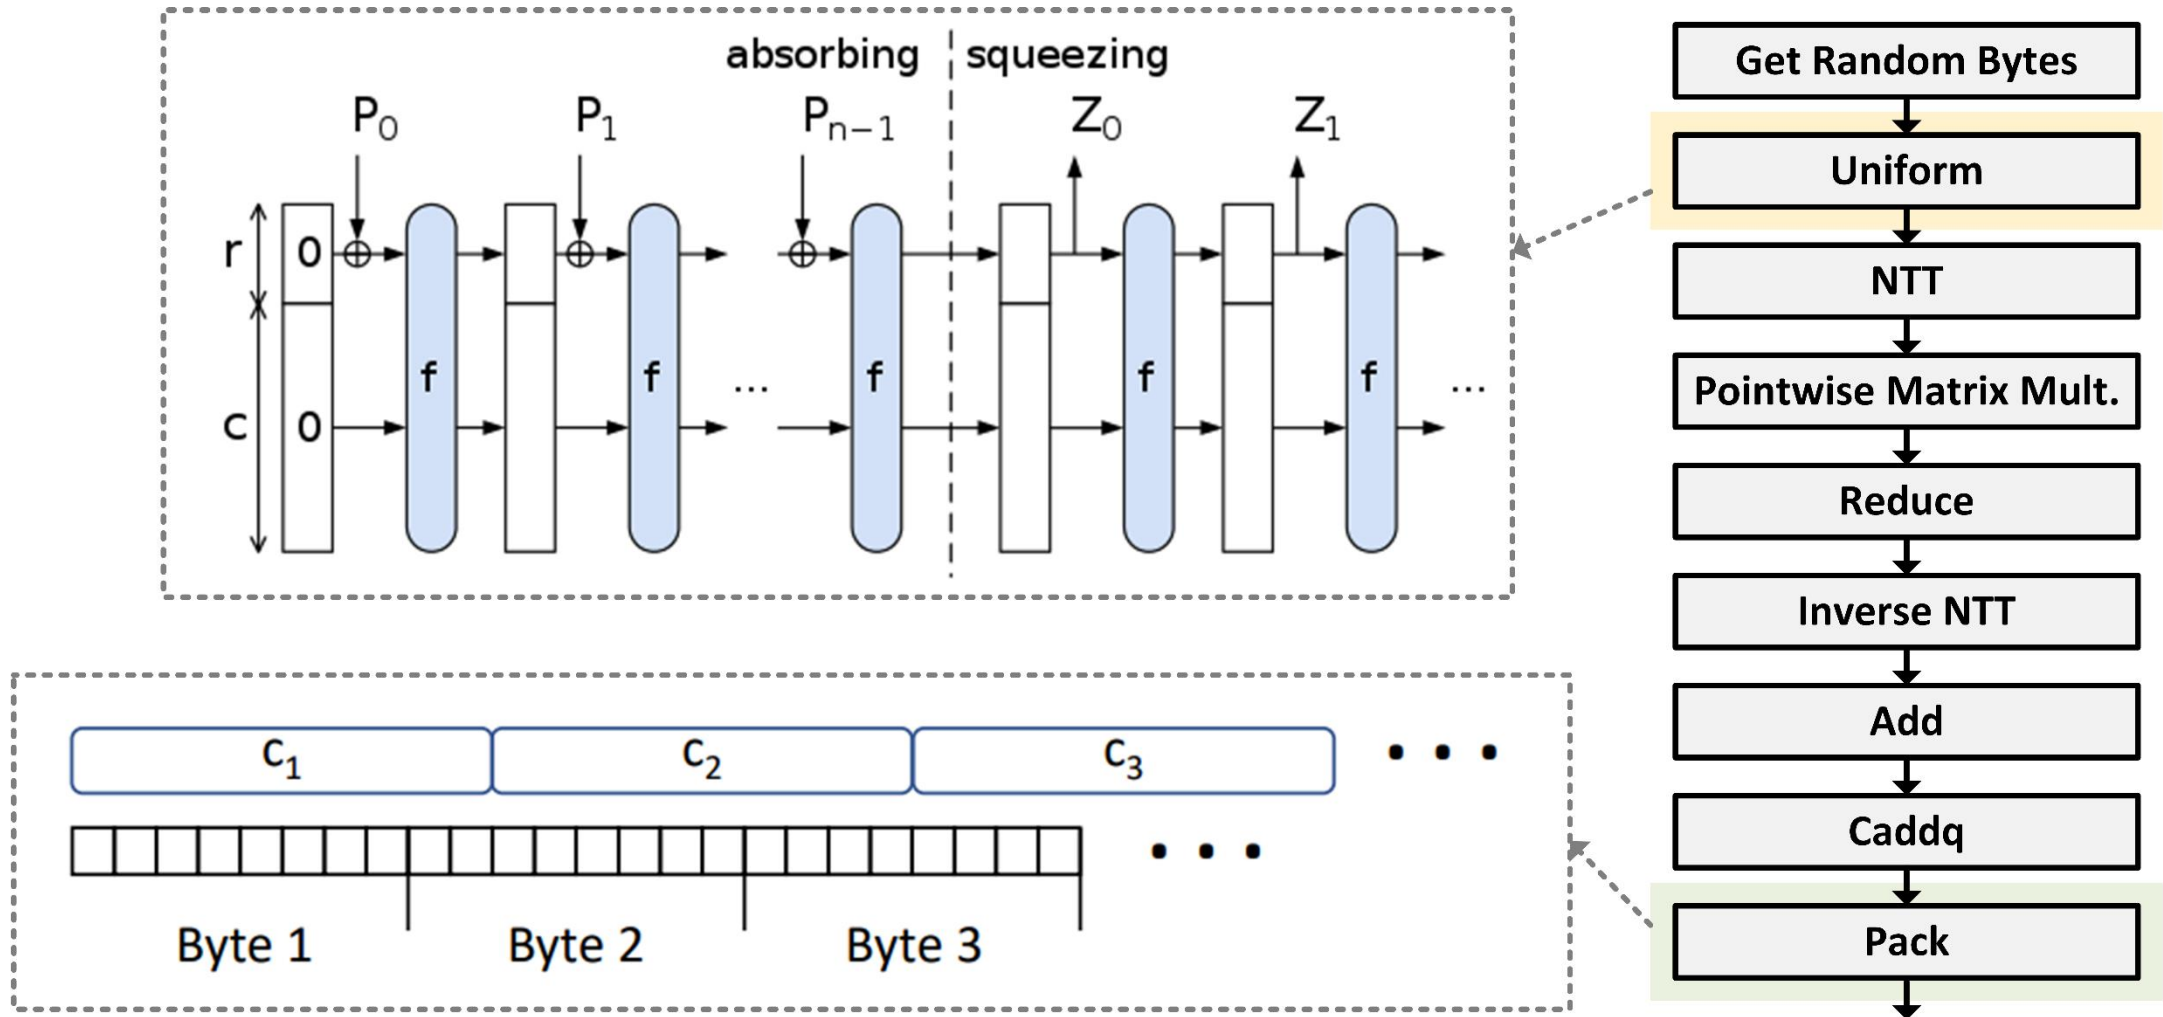


Proposed Architecture

- Fully-Integrated architecture supporting operations for Dilithium

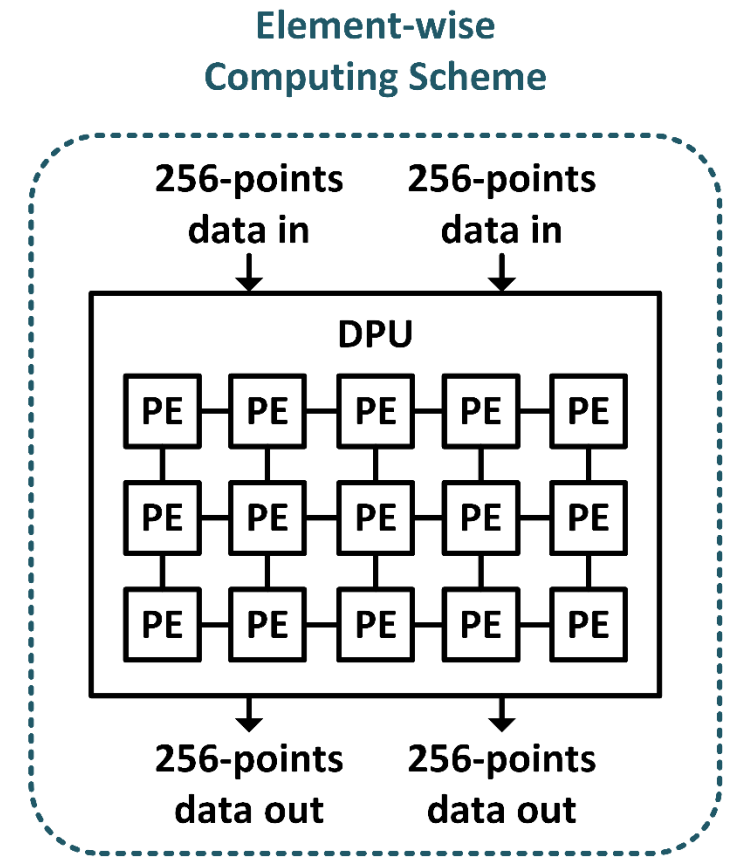
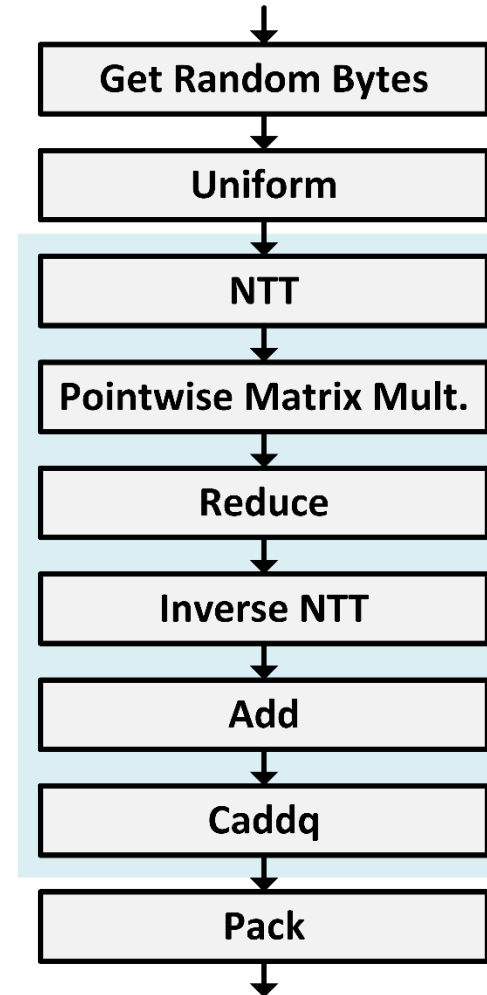


SPU and Pack Unit



DPU

- Support operations for element-wise data computation
- 256-points data is computed in a parallel manner
- The DPU is composed of 256 PEs
- Each processing element (PE) supports different operation
- Hardware resource sharing across similar operations

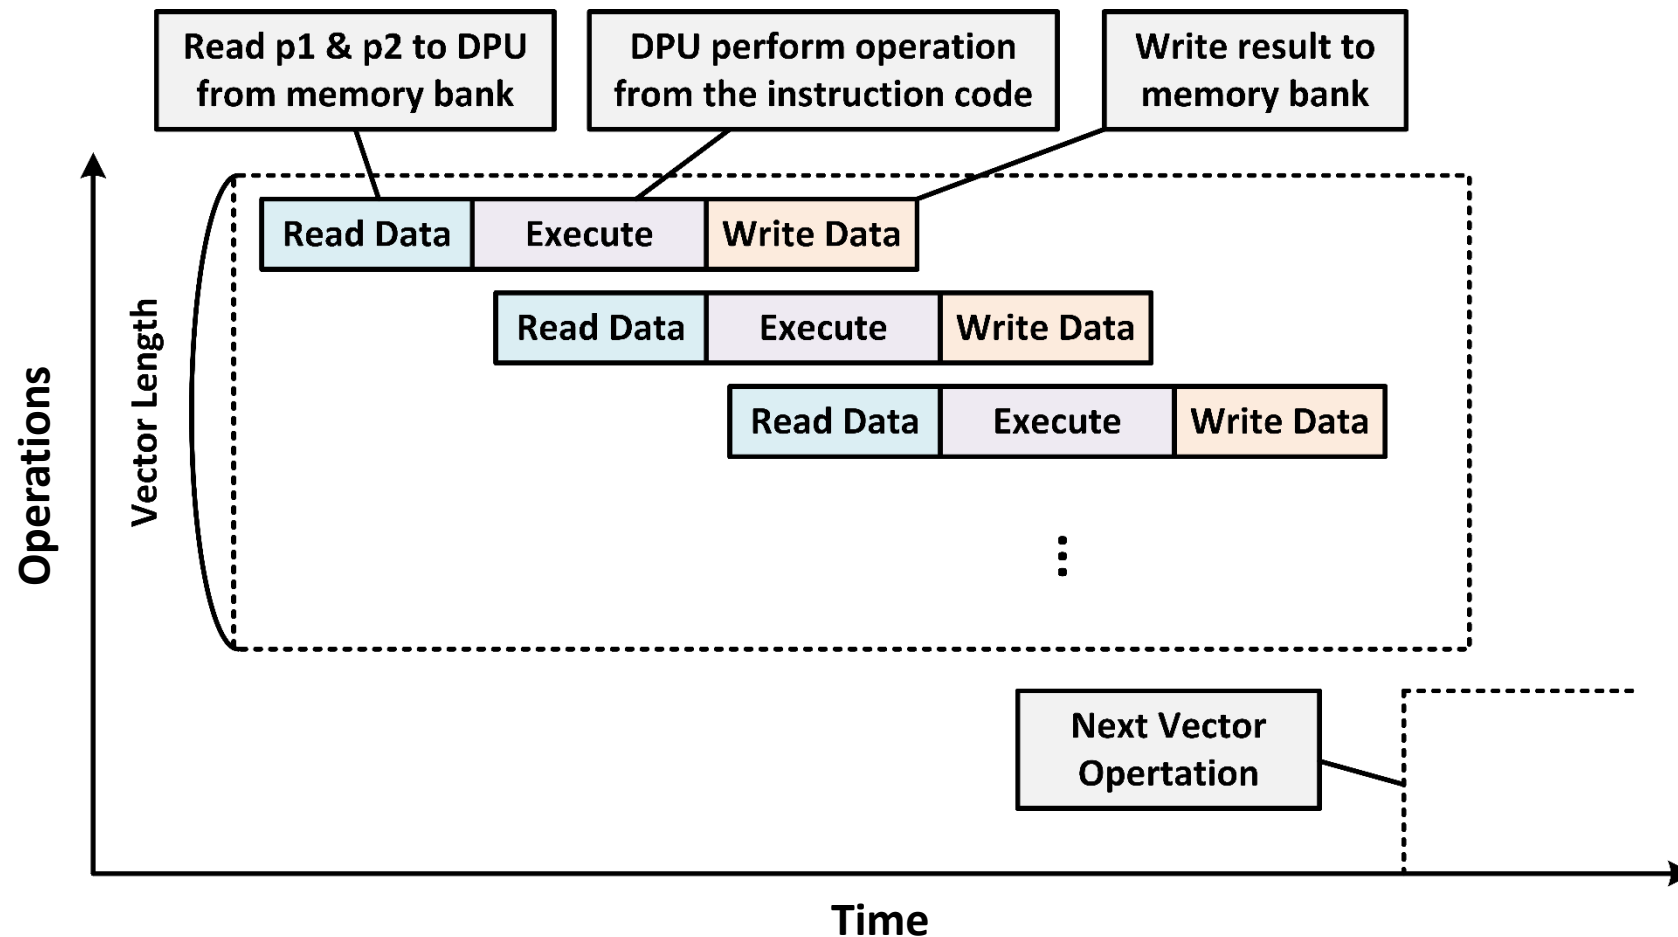


Supported PE Operations

- Addition
- Subtraction
- Multiplication
- CADDQ: Add Q if the data is smaller than 0
- Reduce: Reduce the range of data to $0 \sim Q$
- Pass
- Power2Round: Decompose a data into MSBs and LSBs
- NTT
- Inverse NTT

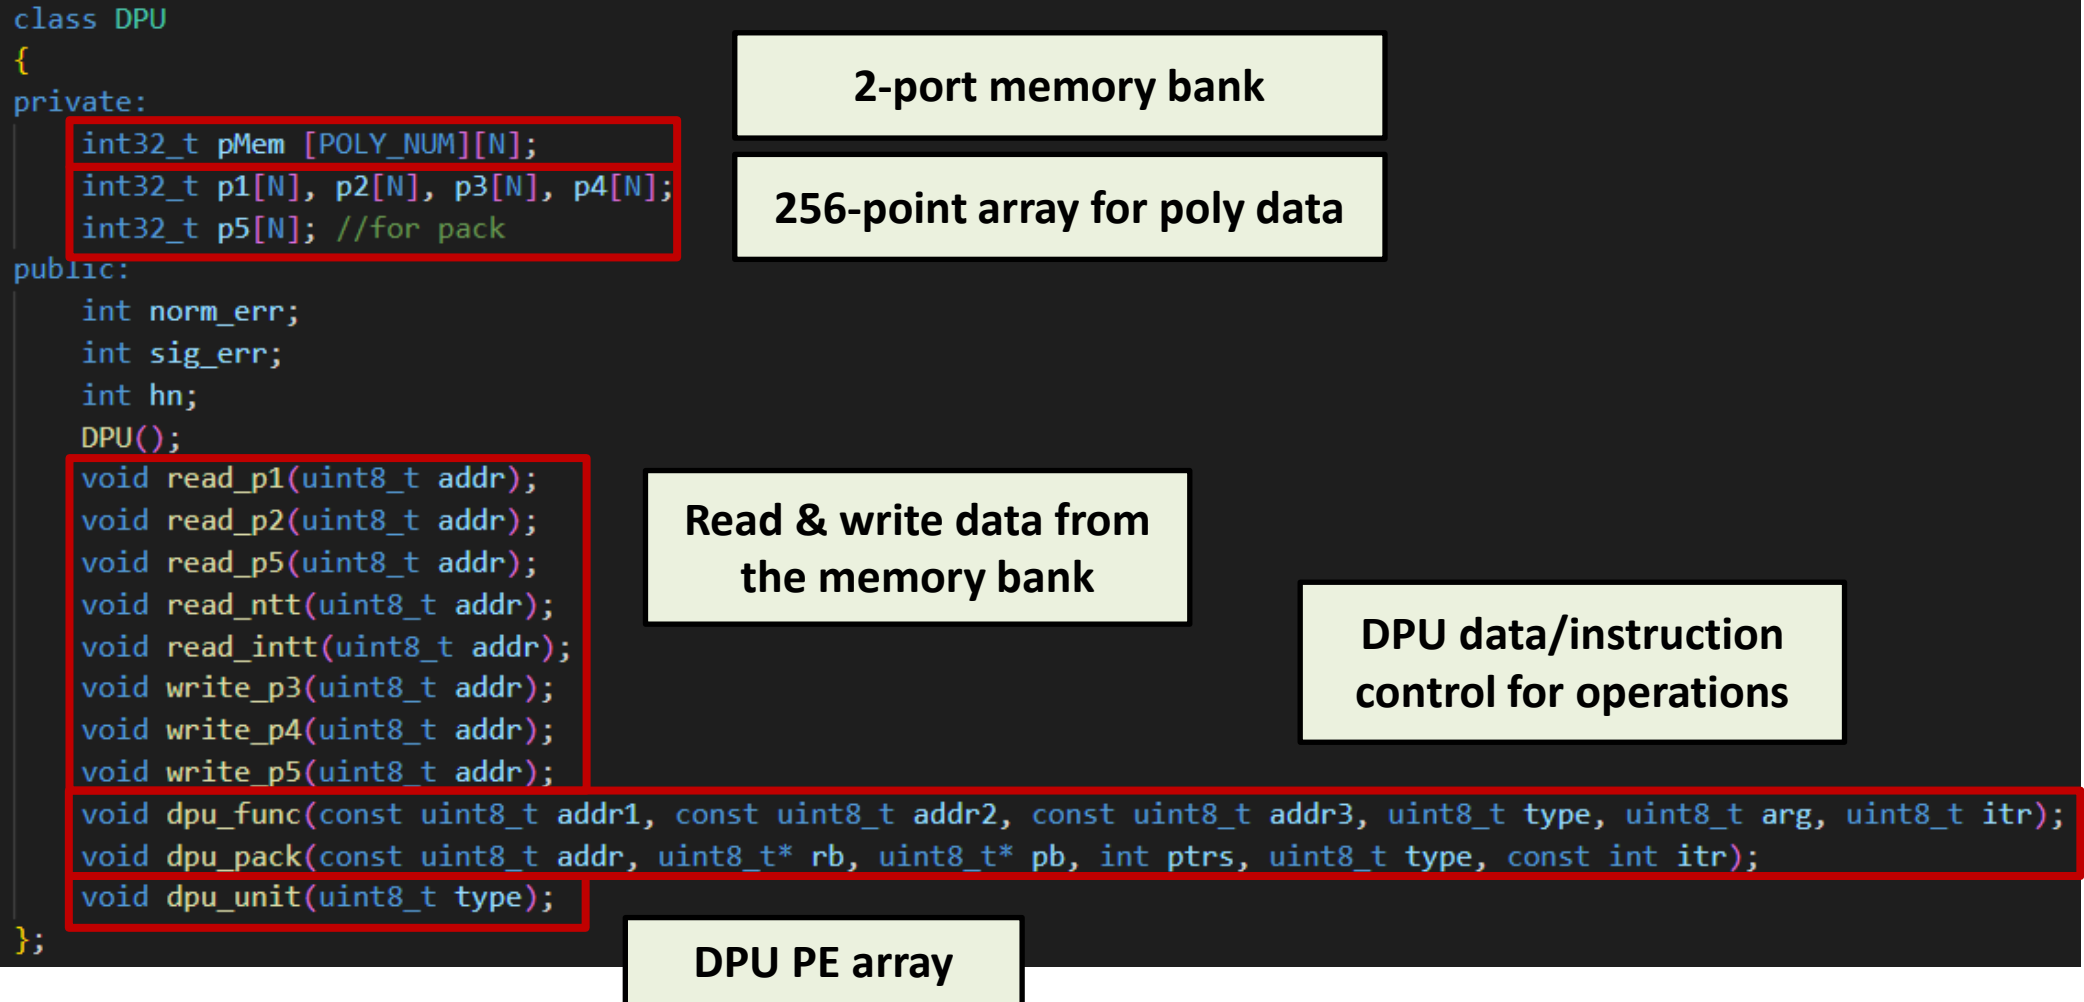
DPU Operation Flow

- Pipelined operation for DPU unit
 - Fully utilize all 256 PEs in DPU



Code Implementation – DPU unit

Code Structure for DPU Class



Code Implementation – DPU unit

Code Structure for DPU Modules

- Unroll loops to facilitate parallel processing

```
void DPU::read_p1(uint8_t addr){  
    for(int i=0;i<N;i++) {  
        #pragma HLS unroll  
        p1[i] = pMem[addr][i];  
    }  
}
```

```
void DPU::write_p3(uint8_t addr){  
    for(int i=0;i<N;i++) {  
        #pragma HLS unroll  
        pMem[addr][i] = p3[i];  
    }  
}
```

- Integrate all micro-operations in a single function
- DPU read two inputs and generate two outputs
 - p1 and p2 in
 - p3 and p4 out

```
void DPU::dpu_unit(uint8_t type){  
    static int64_t tmp;  
    for(int i=0;i<N;i++) {  
        #pragma HLS UNROLL  
        switch(type){  
            case OP_ADD:  
                p3[i] = p1[i] + p2[i];  
                break;  
            case OP_SUB:  
                p3[i] = p1[i] - p2[i];  
                break;  
            case OP_MUL:  
                tmp = (int64_t)p1[i]*p2[i];  
                p3[i] = tmp;  
                p4[i] = (int32_t)(tmp>>32);  
                break;  
            case OP_RD32:  
                p3[i] = (p1[i] + (1 << 22)) >> 23;  
                break;  
            case OP_CADDQ:  
                p3[i] = p1[i] + ((p1[i] >> 31) & Q);  
                break;  
            case OP_PASS:  
                p3[i] = p1[i];  
                break;  
            case OP_POW2ROUND://p1[i](a) --> p3[i](a1), p4[i](a2)  
                p3[i] = (p1[i] + (1 << (D-1)) - 1) >> D;  
                p4[i] = p1[i] - (p3[i] << D);  
                break;  
            case OP NTT:
```

Code Implementation – DPU unit

Code Structure for DPU Operations

- For each function, an array of poly can be processed independently
- No data dependency between each array element

Prevent the modules from over-duplicating

Set false dependency for loop pipelining

3 operations (Read, Operation, Write) for each for loop

```
void DPU::dpu_func(const uint8_t addr1, const uint8_t addr2, const uint8_t addr3, const uint8_t itr) {
    #pragma HLS allocation function instances=dpu_unit limit=1
    #pragma HLS allocation function instances=read_poly limit=1
    #pragma HLS allocation function instances=write_poly limit=1

    uint8_t tmp_addr;
    uint8_t tmp1_addr, tmp2_addr;

    switch(type){
        case FUNC_ADD:
            dpu_add: for(int i=0;i<itr;i++){
                // #pragma HLS pipeline
                #pragma HLS DEPENDENCE variable=pMem1 inter false
                #pragma HLS DEPENDENCE variable=pMem2 inter false
                #pragma HLS DEPENDENCE variable=p1 inter false
                #pragma HLS DEPENDENCE variable=p2 inter false
                #pragma HLS DEPENDENCE variable=p3 inter false
                #pragma HLS DEPENDENCE variable=p4 inter false

                read_poly(addr1+i, addr2+i, true, true);
                dpu_unit(OP_ADD);
                write_poly(addr3+i, 0, true, false);
            }
            break;


        case FUNC_SUB:
            dpu_sub: for(int i=0;i<itr;i++){
                #pragma HLS DEPENDENCE variable=pMem1 inter false
                #pragma HLS DEPENDENCE variable=pMem2 inter false
                #pragma HLS DEPENDENCE variable=p1 inter false
                #pragma HLS DEPENDENCE variable=p2 inter false
                #pragma HLS DEPENDENCE variable=p3 inter false
```

Code Implementation – DPU unit

Code Structure for Top Level Function

Limiting the number of
DPU modules (to 1)

Array reshape for efficient memory allocation




Call DPU function for
each operations
(Instruction code pass-in)

```
#pragma HLS allocation function instances=dpu.dpu_func limit=1
#pragma HLS ARRAY_RESHAPE variable=dpu.pMem type=complete dim=2
#pragma HLS bind_storage variable=dpu.pMem type=RAM_T2P
#pragma HLS ARRAY_RESHAPE variable=dpu.p1 type=complete dim=1
#pragma HLS ARRAY_RESHAPE variable=dpu.p2 type=complete dim=1
#pragma HLS ARRAY_RESHAPE variable=dpu.p3 type=complete dim=1
#pragma HLS ARRAY_RESHAPE variable=dpu.p4 type=complete dim=1

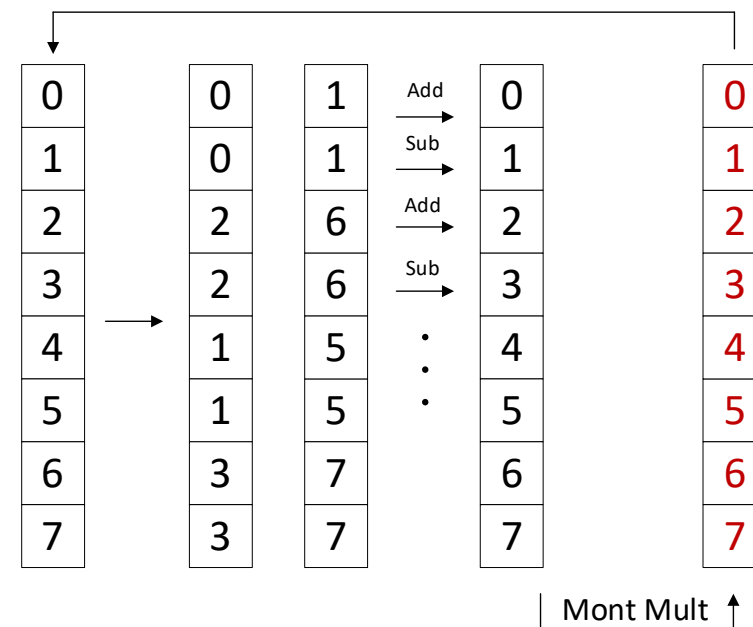
ntt:
    dpu.dpu_func( S1H_ADDR, 0, 0, FUNC_NTT, 0, L);
matmul:
    dpu.dpu_func( A_ADDR, S1H_ADDR, T1_ADDR, FUNC_MATMUL, 0, L);
rd:
    dpu.dpu_func( T1_ADDR, 0, 0, FUNC_RD, 0, K);
intt:
    dpu.dpu_func( T1_ADDR, 0, 0, FUNC_INTT, 0, K);
montmul:
    dpu.dpu_func(MONT2_ADDR, T1_ADDR, T1_ADDR, FUNC_MONTMUL, 0, K);
add:
    dpu.dpu_func( T1_ADDR, S2_ADDR, T1_ADDR, FUNC_ADD, 0, K);
caddq:
    dpu.dpu_func( T1_ADDR, 0, 0, FUNC_CADDQ, 0, K);
pow2round:
    dpu.dpu_func( T1_ADDR, T1_ADDR, T0_ADDR, FUNC_POW2ROUND, 0, K);
pack1:
    dpu.dpu_pack( 0, rho, pk, 0, BYTE_PACK, SEEDBYTES);
pack2:
    dpu.dpu_pack(T1_ADDR, 0, pk, SEEDBYTES, T1_PACK, K);
    spu.shake(256, tr, CRHBYTES, pk, CRYPTO_PUBLICKEYBYTES);
pack3:
```

Code Implementation – NTT and Inverse NTT

- NTT
 - Step1: permutation
 - Step3: vector multiplication
 - Step2: vector add/sub
 - Step4: go to step1



- Inverse NTT
 - Step1: permutation
 - Step2: vector add/sub
 - Step3: vector multiplication
 - Step4: go to step1



Code Implementation – SPU unit and Pack unit

- SPU & pack unit

```
class SPU
{
private:
    uint64_t s[25];
    unsigned int pos;
public:
    SPU();
    void store64(uint8_t x[8], uint64_t u);
    uint64_t load64(const uint8_t x[8]);
    void KeccakF1600_StatePermute();
    void shake_init();
    void stream_init(unsigned int mode, const uint8_t seed[CRHBYTES], uint16_t n);
    void shake_absorb(unsigned int mode, const uint8_t *m, size_t mlen);
    void shake_finalize(unsigned int mode);
    void shake_squeezeblocks(unsigned int mode, uint8_t *out, size_t nblocks);
    void shake_squeeze(unsigned int mode, uint8_t *out, size_t outlen);
    void shake(unsigned int mode, uint8_t *out, size_t outlen, const uint8_t *in);
    void sample_uniform(int32_t* a, const uint8_t seed[SEEDBYTES], uint16_t nonce);
    void sample_eta(int32_t* a, const uint8_t seed[SEEDBYTES], uint16_t nonce);
    void sample_cp(int32_t *c, const uint8_t seed[SEEDBYTES]);
};
```

```
void DPU::dpu_pack(const uint8_t addr, uint8_t* rb, uint8_t* pb, int ptrs, uint8_t type, const int itr){
#pragma HLS ARRAY_RESHAPE variable=p5 type=complete dim=1
    static int ptr;
    ptr = ptrs;
    switch(type){
        case T1_PACK:
            for(int j=0; j<K; j++){
                read_p5(addr+j);
                for(int i=0; i < N/4; ++i) {
                    pb[ptr+0] = (p5[4*i+0] >> 0);
                    pb[ptr+1] = (p5[4*i+0] >> 8) | (p5[4*i+1] << 2);
                    pb[ptr+2] = (p5[4*i+1] >> 6) | (p5[4*i+2] << 4);
                    pb[ptr+3] = (p5[4*i+2] >> 4) | (p5[4*i+3] << 6);
                    pb[ptr+4] = (p5[4*i+3] >> 2);
                    ptr += 5;
                }
            }
            break;
        case T1_UNPACK:
            for(int j=0; j<K; j++){
                for(int i = 0; i < N/4; ++i) {
                    p5[4*i+0] = (((uint32_t)pb[ptr+0] >> 0) | ((uint32_t)pb[ptr+1] << 8)) & 0x3FF;
                    p5[4*i+1] = (((uint32_t)pb[ptr+1] >> 2) | ((uint32_t)pb[ptr+2] << 6)) & 0x3FF;
                    p5[4*i+2] = (((uint32_t)pb[ptr+2] >> 4) | ((uint32_t)pb[ptr+3] << 4)) & 0x3FF;
                    p5[4*i+3] = (((uint32_t)pb[ptr+3] >> 6) | ((uint32_t)pb[ptr+4] << 2)) & 0x3FF;
                    ptr += 5;
                }
                write_p5(addr+j);
            }
            break;
        case T0_PACK:
            for(int j=0; j<K; j++){
                read_p5(addr+j);
                for(int i = 0; i < N/8; ++i) {
```

HLS Observations

1. DPU Unit high latency and recourse usage
 - Use “if-else” instead of “switch-case”

Latency (cycles)		Latency (absolute)		Interval		Pipeline
min	max	min	max	min	max	Type
64	1346	0.640 us	13.460 us	64	1346	no

Name	BRAM_18K	DSP	FF	LUT
DSP	-	-	-	-
Expression	-	-	0	118
FIFO	-	-	-	-
Instance	-	1536	1819846	974722
Memory	-	-	-	-
Multiplexer	-	-	-	3163
Register	-	-	172443	-
Total	0	1536	1992289	978003
Available SLR	1344	3072	864000	432000
Utilization SLR (%)	0	50	230	226
Available	4032	9216	2592000	1296000
Utilization (%)	0	16	76	75

Latency (cycles)		Latency (absolute)		Interval		Pipeline
min	max	min	max	min	max	Type
1	7803	10.000 ns	78.030 us	1	7803	no

Name	BRAM_18K	DSP	FF	LUT	URAM
Expression	-	-	0	126	-
Instance	-	1536	1303268	237909	-
Memory	-	-	-	-	-
Multiplexer	-	-	-	3062	-
Register	-	-	270677	-	-
Total	0	1536	1573945	241097	0
Available SLR	1344	3072	864000	432000	320
Utilization SLR (%)	0	50	182	55	0
Available	4032	9216	2592000	1296000	960
Utilization (%)	0	16	60	18	0

```
void DPU::dpu_func(const uint8_t addr1,
#pragma HLS allocation function instance
#pragma HLS allocation function instance
#pragma HLS allocation function instance

uint8_t tmp_addr;
uint8_t tmp1_addr, tmp2_addr;

switch (type)
case FUNC_ADD:
dpu_add: for(int i=0;i<itr;i++){
// #pragma HLS pipeline
#pragma HLS DEPENDENCE variable=pMem1 in
#pragma HLS DEPENDENCE variable=pMem2 in
#pragma HLS DEPENDENCE variable=p1 inter
#pragma HLS DEPENDENCE variable=p2 inter
#pragma HLS DEPENDENCE variable=p3 inter
#pragma HLS DEPENDENCE variable=p4 inter
    read_poly(addr1+i, addr2
    dpu_unit(OP_ADD);
    write_poly(addr3+i, 0, t
    }
    break;

case FUNC_SUB:
dpu_sub: for(int i=0;i<itr;i++){
#pragma HLS DEPENDENCE variable=pMem1 in
#pragma HLS DEPENDENCE variable=pMem2 in
#pragma HLS DEPENDENCE variable=p1 inter
```

HLS Observations

2. Use ALLOCATION pragma to restrict instance number

```
#pragma HLS allocation function instances=dpu_unit limit=1
#pragma HLS allocation function instances=read_poly limit=1
#pragma HLS allocation function instances=write_poly limit=1
```

```
void DPU::dpu_func(const uint8_t addr1,
#pragma HLS allocation function instance
#pragma HLS allocation function instance
#pragma HLS allocation function instance
```

3. Inline read/write function to set false dependency

- Vitis HLS cannot set false dependency for member in function
- Add inline pragma in read/write function

```
uint8_t tmp_addr;
uint8_t tmp1_addr, tmp2_addr;

switch(type){
    case FUNC_ADD:
dpu_add:   for(int i=0;i<itr;i++){
// #pragma HLS pipeline
#pragma HLS DEPENDENCE variable=pMem1 in
#pragma HLS DEPENDENCE variable=pMem2 in
#pragma HLS DEPENDENCE variable=p1 inter
#pragma HLS DEPENDENCE variable=p2 inter
#pragma HLS DEPENDENCE variable=p3 inter
#pragma HLS DEPENDENCE variable=p4 inter
                read_poly(addr1+i, addr2
                dpu_unit(OP_ADD);
                write_poly(addr3+i, 0, t
            }
            break;

    case FUNC_SUB:
dpu_sub:   for(int i=0;i<itr;i++){
#pragma HLS DEPENDENCE variable=pMem1 in
#pragma HLS DEPENDENCE variable=pMem2 in
#pragma HLS DEPENDENCE variable=p1 inter
```

HLS Implementation Result

- Total execution time: 87646
 - SPU and Pack Unit is the bottleneck

▶ dpu_func	271	660	28	238	658	26
▶ dpu_func_Pipeline_FUNC_POW2ROUND_LOOP1				14	14	14
▶ dpu_func_Pipeline_FUNC_CADDQ_LOOP1				13	13	13
▶ dpu_func_Pipeline_FUNC_MONTMUL_LOOP1				13	13	13
▶ dpu_func_Pipeline_FUNC_RD_LOOP1				13	13	13
▶ dpu_func_Pipeline_FUNC_ADD_LOOP1				13	13	13
▶ dpu_func_Pipeline_FUNC_MONTMUL_LOOP2				13	13	13
▶ dpu_func_Pipeline_FUNC_MONTMUL_LOOP3				13	13	13
▶ dpu_func_Pipeline_FUNC_MONTMUL_LOOP4				13	13	13
▶ dpu_func_Pipeline_FUNC_RD_LOOP2				13	13	13
▶ dpu_func_Pipeline_FUNC_RD_LOOP3				13	13	13
▶ FUNC_MATMUL_LOOP0				475	475	475
▶ FUNC_NTT_LOOP0				577	577	577
▶ FUNC_INTT_LOOP0				657	657	657

Modules & Loops	Avg II	Max II	Min II	Avg Latency	Max Latency	Min Latency
▶ dpu_keygen				87646	87646	87646
▶ dpu_keygen_Pipeline_VITIS_LOOP_40_1				256	256	256
▶ dpu_keygen_Pipeline_VITIS_LOOP_41_2				256	256	256
▶ dpu_keygen_Pipeline_VITIS_LOOP_321_1				25	25	25
▶ dpu_keygen_Pipeline_VITIS_LOOP_42_3_VITIS_LOOP_43_4				2048	2048	2048
▶ shake_absorb_1	1945	4390	1860	28	28	28
▶ dpu_keygen_Pipeline_VITIS_LOOP_48_5_VITIS_LOOP_49_6				2048	2048	2048
▶ dpu_keygen_Pipeline_VITIS_LOOP_54_7				257	257	257
▶ dpu_keygen_Pipeline_VITIS_LOOP_55_8				257	257	257
▶ dpu_keygen_Pipeline_VITIS_LOOP_56_9				257	257	257
▶ shake_squeeze_2				101	101	101
▶ dpu_func	271	660	28	238	658	26
▶ dpu_keygen_Pipeline_VITIS_LOOP_55_5				32	32	32
▶ dpu_keygen_Pipeline_VITIS_LOOP_321_110				25	25	25
▶ dpu_keygen_Pipeline_VITIS_LOOP_60_6				32	32	32
▶ dpu_keygen_Pipeline_VITIS_LOOP_61_7				32	32	32
▶ dpu_pack				2725	2725	2725
▶ shake_absorb_3				1734	1734	1734
▶ shake_squeeze				77	77	77
▶ dpu_keygen_Pipeline_VITIS_LOOP_62_8				48	48	48
▶ dpu_pack_4	2149	2344	1954	2404	2918	1952
▶ VITIS_LOOP_31_1_VITIS_LOOP_32_2				55828	55828	55828
▶ VITIS_LOOP_37_3				6791	6791	6791
▶ VITIS_LOOP_42_4				6669	6669	6669

Design Comparison

- Compared with high-end CPU & state-of-the-art FPGA architecture
 - CPU: Intel Core i7-9700
 - FPGA Implementation: ICFPT'21[3]

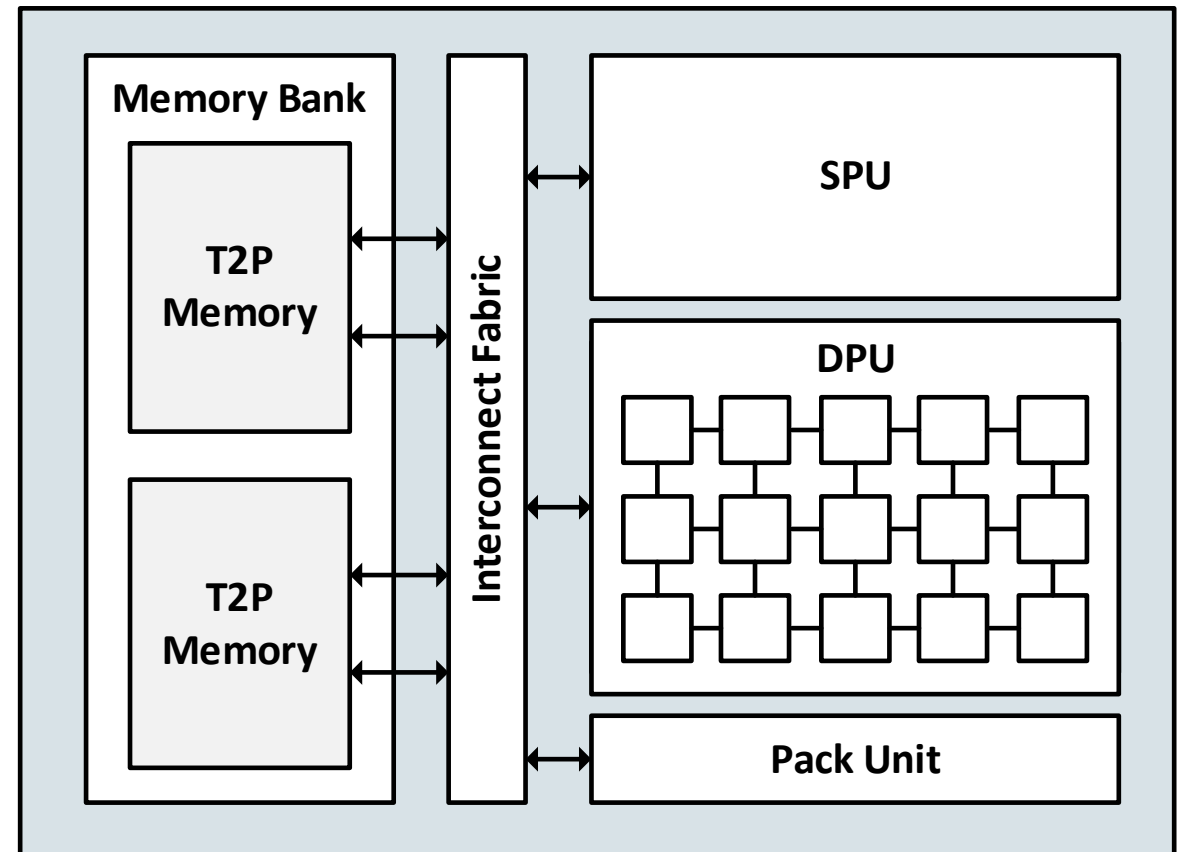
Latency (cycle)

	CPU	This Work
Total Key Generation	446475	87646

	ICFPT'21[3]	This Work
NTT	1831	577
Matrix Multiplication	1754	475
Inverse NTT	1475	670

Architecture Exploration

- Read p1 and read p2 takes two cycle
 - Additional data idle time can increase buffer usage
 - Same situation happens for memory write access
 - **2 memory banks**
 - Complicated data access pattern for different operations
 - How to allocate data?
- ➔ **Algorithm breakdown & analysis**
- ➔ **Intermediate product reallocation**



Architecture Exploration

2-memory-bank implementation

- Increase memory bandwidth via efficient allocation of variable
 - Enhance pipeline efficiency
 - Read & write in one cycle

```
switch(type){
    case FUNC_ADD:
dpu_add:   for(int i=0;i<itr;i++){
// #pragma HLS pipeline
#pragma HLS DEPENDENCE variable=
#pragma HLS DEPENDENCE variable=
#pragma HLS DEPENDENCE variable=
#pragma HLS DEPENDENCE variable=
#pragma HLS DEPENDENCE variable=p3 inter false
#pragma HLS DEPENDENCE variable=p4 inter false
        read_poly(addr1+i, addr2+i, true, true);
        dpu_unit(OP_ADD);
        write_poly(addr3+i, 0, true, false);
    }
    break;
```

Clearly specify the three operations

```
void DPU::read_poly(uint8_t addr_p1, uint8_t addr_p2, bool p1_en, bool p2_en)
#ifdef DEBUG
    assert((!(p1_en && p2_en)) || ((addr_p1 >= (1 << 7)) ^ (addr_p2 >= (1 << 7))))
#endif
    bool mem_switch;
    if ((!isInMem1(addr_p1) && p1_en) || (isInMem1(addr_p2) && p2_en))
        else mem_switch = 0;

    uint8_t readMem1_addr, readMem2_addr;
    if (mem_switch) {
        readMem1_addr = addr_p2;
        readMem2_addr = addr_p1;
    }
    else {
        readMem1_addr = addr_p1;
        readMem2_addr = addr_p2;
    }

    if (readMem2_addr > 127) readMem2_addr = getMem2Addr(readMem2_addr);

    for(int i=0;i<N;i++) {
#pragma HLS UNROLL
        if (mem_switch) {
            p1[i] = pMem2[readMem2_addr][i];
            p2[i] = pMem1[readMem1_addr][i];
        }
        else {
            p1[i] = pMem1[readMem1_addr][i];
            p2[i] = pMem2[readMem2_addr][i];
        }
    }
}
```

Assertion for ZERO memory conflict

Memory address & data switching

Architecture Exploration

Encountered Problem

- Excessively large resource usage for data switching in read module
- Generate numerous expressions using LUT

select_ln107_1002_fu_23257_p3	select	0	0	4095	1	8192
select_ln107_1003_fu_23263_p3	select	0	0	4095	1	8192
select_ln107_1004_fu_23292_p3	select	0	0	4095	1	8192
select_ln107_1005_fu_23298_p3	select	0	0	4095	1	8192
select_ln107_1006_fu_23325_p3	select	0	0	4095	1	8192
select_ln107_1007_fu_23331_p3	select	0	0	4095	1	8192
select_ln107_1008_fu_23358_p3	select	0	0	4095	1	8192
select_ln107_1009_fu_23364_p3	select	0	0	4095	1	8192
select_ln107_100_fu_8047_p3	select	0	0	4095	1	8192
select_ln107_1010_fu_23382_p3	select	0	0	4095	1	8192
select_ln107_1011_fu_23415_p3	select	0	0	4095	1	8192
select_ln107_1012_fu_23399_p3	select	0	0	4095	1	8192
select_ln107_1013_fu_23431_p3	select	0	0	4095	1	8192
select_ln107_1014_fu_23459_p3	select	0	0	4095	1	8192
select_ln107_1015_fu_23465_p3	select	0	0	4095	1	8192
select_ln107_1016_fu_23494_p3	select	0	0	4095	1	8192
select_ln107_1017_fu_23500_p3	select	0	0	4095	1	8192
select_ln107_1018_fu_23527_p3	select	0	0	4095	1	8192
select_ln107_1019_fu_23533_p3	select	0	0	4095	1	8192
select_ln107_101_fu_8079_p3	select	0	0	4095	1	8192
select_ln107_1020_fu_23560_p3	select	0	0	4095	1	8192
select_ln107_1021_fu_23566_p3	select	0	0	4095	1	8192

```
=====
== Utilization Estimates
=====
* Summary:
```

Name	BRAM_18K	DSP	FF	LUT	URAM
DSP	-	-	-	-	-
Expression	-	-	0	4193310	-
FIFO	-	-	-	-	-
Instance	-	-	-	-	-
Memory	-	-	-	-	-
Multiplexer	-	-	-	-	-
Register	-	-	4243441	16288	-
Total	0	0	4243441	4209598	0
Available SLR	1344	3072	864000	432000	320
Utilization SLR (%)	0	0	491	974	0
Available	4032	9216	2592000	1296000	960
Utilization (%)	0	0	163	324	0

Summary & Conclusions

- **Fully-Integrated architecture for Dilithium using high level synthesis**
 - Approximately good as state-of-the-art design

Difficulties

- Rewriting the C code takes time
 - 2000+ lines code rewrite
 - 2-3 weeks labor work for 2 people

What's next? (If possible)

- Finish two-memory bank design
- Manual resource sharing in the DPU PE array
- Optimization on SPU and Pack Unit

References

1. NIST Post-Quantum Cryptography Standardization: <https://csrc.nist.gov/Projects/post-quantum-cryptography>
2. Dilithium: <https://pq-crystals.org/dilithium/index.shtml>
3. L. Beckwith, D. T. Nguyen and K. Gaj, "High-Performance Hardware Implementation of CRYSTALS-Dilithium," *2021 International Conference on Field-Programmable Technology (ICFPT)*, pp. 1-10, 2021.